ANKIT KUMAR CHOUDHARY

LLM Coding-Agent Evaluation Engineer | Backend & Distributed Systems (SDE-2)
+91-7067423108 | ankitchoudhary.sde@gmail.com | LinkedIn | Portfolio | Bengaluru, India (Remote)

SUMMARY

Backend engineer (Goldman Sachs, Spinny) who builds and grades evaluation tasks for frontier LLM coding agents. I author SWE-bench-style tasks on real open-source codebases: reproducible Docker environments, hidden fail-to-pass and regression tests, and rubrics that catch reward hacking and fabricated claims. 3 years shipping Java/Spring Boot and Python/Django microservices at 120 TPS and 70 ms P99. Codeforces Expert (1800+).

EXPERIENCE

LLM Coding-Agent Evaluation Engineer (Contract) · Freelance · Remote 10/2026 – Present

- **Task authoring:** Authored ~30 verified SWE-bench-style tasks on real open-source Python and Java codebases (up to 355k lines), validated across ~70 agent trials.
- **Reproducible environments:** Built Docker sandboxes with pinned commits, PostgreSQL and Redis; every environment builds in under 15 minutes.
- **Deterministic grading:** Wrote hidden fail-to-pass and regression tests (pytest), each verified three ways: starting code fails, reference solution passes, plausible wrong fixes fail.
- **Distributed-systems tasks:** Ported a 33-commit Redis Cluster PR onto a newer master and graded it on a 3-node Docker cluster; built lock-ownership, shutdown-race and monotonic-clock migration tasks.
- **Anti-reward-hacking:** Hardened graders so a skip-every-test exploit that passed the old grader now fails; fixed OpenBLAS oversubscription in emulated containers, cutting test time from 7.5 s to 0.45 s.
- **Trajectory review and QA:** Graded agent trajectories on 8-dimension rubrics with fabrication checks; as QA reviewer, caught errors in 7 of 12 rubrics on a peer submission and wrote a source-verified review checklist.

Career Break 05/2025 – 10/2026

- Took planned time off for a personal health matter, now back at work full-time.
- Upskilled in LLM coding agents, agent evaluation tooling and system design, which led directly to my current evaluation contract work.

Software Developer II (SDE-2) · Spinny · Gurugram 12/2024 – 05/2025

- **Anomaly detection:** Designed a distributed pipeline over call and activity logs from 2,500+ field inspectors (Python, Django, MongoDB, Redis, Firebase), with Redis deduplication and real-time fraud alerts.
- **Impact:** Flagged 300+ anomalies a month, cut lead leakage by 22% and saved ₹1.3 Cr (~$150K) per quarter.
- **Cost prediction:** Built a parts and labor cost engine for 5,000+ daily inspections using a Camunda BPM issue-probability matrix, cutting manual effort by 85%.

Software Developer · Goldman Sachs · Bengaluru 06/2022 – 08/2024

- **Event-driven capture:** Designed a prospect-capture microservice on MongoDB Change Streams (CDC) and Kafka, raising customer engagement by 23%.
- **Performance and security:** Sustained 120 TPS at 70 ms P99 through connection pooling, indexing and @Async, with AES-256 PII encryption and Resilience4j.
- **Distributed rendering:** Scaled a statement-generation service across 10 Hazelcast nodes to 280 PDFs per second.
- **Release platform:** Built an internal APK distribution system on AWS S3 handling 50+ releases a month with zero downtime and 200 ms version checks.

SKILLS

AI evaluation: LLM coding-agent evaluation, SWE-bench-style task authoring, RLHF data, agent trajectory analysis, rubric design, reward-hacking detection, code review, prompt design, Harbor, OpenCode, mini-swe-agent

Languages: Python, Java, C++, SQL

Backend: Spring Boot, Django, Flask, microservices, REST APIs, event-driven architecture, concurrency, system design

Data and infrastructure: Docker, PostgreSQL, MongoDB, Redis, Kafka, Hazelcast, Firebase, AWS S3, Git, pytest

EDUCATION

B.Tech, Computer Science and Engineering · NIT Patna 2018 – 2022